


Computer Vision

Fundamentals of Artificial Intelligence

Goal of Computer Vision

extract semantic information from digital image data
to be used for decision making support or automated systems

challenging problem:
images are only 2D projections
of the 3D world



nowadays heavily powered by artificial intelligence (AI), especially machine learning (ML)

Applications of Computer Vision

facial recognition



automated inspection



autonomous driving



medical imaging



optical character recognition




augmented reality



and many more ...

Digital Image Acquisition Process





Photons to Electrons

largely replaced chemical and analog imaging




single sensor: photodiode

sensor strips: motion perpendicular to strip for imaging in other direction, used in CT



sensor arrays: used in digital cameras (in focal plane, outputs proportional to integral of light received at each sensor)



Digitizing: Sampling & Quantization



digitizing the coordinate
values: pixels



digitizing the amplitude values:
discrete intensity levels

Discrete Intensity Levels

image data: 2D grid (matrix) of pixels with different intensity values


intensity resolution:

common to use one byte (2^8) to express possible intensity values (8-bit image)

→ 0 = black, 255 = white

storage requirement for, e.g., 1024×1024 8-bit image (grayscale): ~1 MB

(× 3 for RGB color images)





[source](#)

Image as Two-Dimensional Function

image can be seen as function $f(x, y)$:
intensity at position (x, y)

digital image: discrete (sampled and
quantized) version of this function

$$f(x, y) = \begin{bmatrix} f(0, 0) & f(0, 1) & \cdots & f(0, N - 1) \\ f(1, 0) & f(1, 1) & \cdots & f(1, N - 1) \\ \vdots & \vdots & & \vdots \\ f(M - 1, 0) & f(M - 1, 1) & \cdots & f(M - 1, N - 1) \end{bmatrix}$$





Colors in the Human Eye

retina: image sensor of human eye

two kinds of light receptors on surface of retina: rods (monochromatic vision) and cones (color-sensitive cells)

three types of cones: red, green, blue
→ combinations of red, green, and blue create perceived colors



Primary Colors



standardization:

use of specific wavelengths for three primary colors red, green, and blue

→ can reproduce (almost) all visible colors by mixing with different intensities

for printing:

primary colors of pigments each absorb one primary color of light





**PRIMARY AND SECONDARY COLORS
OF LIGHT AND PIGMENT**

RGB Color Model

color model for monitors and cameras
(other color models: CMYK, HSV, ...)


color pixels:



Color Sensing

single image sensor with one of three color filters in front of each pixel

estimation of missing color values for each pixel by interpolation from neighboring pixels
→ full-color image




G	R	G	R
B	G	B	G
G	R	G	R
B	G	B	G


Bayer RGB Pattern

rGb	Rgb	rGb	Rgb
rgB	rGb	rgB	rGb
rGb	Rgb	rGb	Rgb
rgB	rGb	rgB	rGb

Interpolated Pixels



Bayer Pattern
(Color Filter Mosaic)



Raw Image





Interpolated Image

Image Processing

transformations from image to image

(such as scaling, smoothing, sharpening, or contrast stretching)




to facilitate either machine perception or just human interpretation


Image Understanding (Recognition)




Image Classification



Semantic Segmentation



Object Detection



Instance Segmentation

Need for ML

prime example (and also foundation for detection and segmentation):

image classification (whole-image class recognition) according to generic object categories (e.g., cat)

plain keypoint-feature matching only really works for specific instances of a class

→ need to compare with generic objects (e.g., kind of “abstract cat”)

Let’s learn it from data ...




What if you haven’t seen this very cat before?

Image Classification

MNIST

- Modified National Institute of Standards and Technology
- handwritten digits (10 classes)
- black and white images
- 28×28 pixels
- 60k training and 10k test images



CIFAR-10

- Canadian Institute for Advanced Research
- 10 different labeled object classes
- color images
- 32×32 pixels
- 50k training and 10k test images

airplane



automobile



bird



cat



deer



dog



frog



horse



ship




truck




Image Classification with Linear Regression

simplified example: 4-pixel image, 3 classes



need for one common image size per model

Simplified Matrix Multiplication: Bias Trick



geometric interpretation:
separating hyperplanes




matching interpretation:
generic class templates



different rows in matrix W

raw-pixel images not linearly separable
→ linear model has not enough representational power



Cannot separate red and blue points with linear classifier

Image Classification with kNN

choose an image distance, e.g., L1 distance:

$$d(I_1, I_2) = \sum_p |I_1(p) - I_2(p)|$$

test image

56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

training image

10	20	24	17
8	10	89	100
12	16	178	170
4	32	233	112

-

=


pixel-wise absolute value differences

46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108


→ 456

Image Classification with kNN

training examples CIFAR-10 data set



10 nearest neighbors to some test images




better than random guessing, but also not very impressive

In Reality: Lots of Categories



ImageNet

- more than 14 million color images with varying sizes
- more than 20k labeled categories




Recap: Feed-Forward Neural Networks

computation in usual feed-forward network:
matrix multiplication of scalar inputs and weights

$$o_t = \sum_s w_{s,t} i_s$$

dropped dimension of different training observations
(in mini-batch) in this view



Inductive Bias

plain feed-forward networks very inefficient for image classification:
make no use of spatial structure (locality of objects)

→ need to learn it from scratch

better way: introduce it right in the architecture of the ML method
(called inductive bias)

→ convolution with spatial filters (**learned** rather than handcrafted)

Convolutional Neural Networks (CNN or ConvNet)

Grid-Like Data


image data: 2-D grid of pixels (spatial structure)

convolutional networks:

neural networks using learned convolution kernels as weights (in place of general matrix multiplications)

→ highly regularized feed-forward networks

scalar values (like in usual feed-forward network)



[source](#)

Convolution Operation


actually, correlation (convolution would have $-$ here):

$$o_{x,y} = \sum_{s,t} w_{s,t} i_{x+s,y+t}$$

feature map (transformed image) kernel (learned) image (input or feature map)

(again, dropping dimension of different training observations)

2D convolution




Regularization Effects

- sparse interactions: much less weights
- parameter sharing: use same weights for different connections



effect of receptive field over several layers:

- consider only locally restricted number of input values from previous layer
- grows for earlier layers (indirect interactions)
→ hierarchical patterns from simple building blocks (many aspects of nature hierarchical)

convolutional layer:



fully-connected layer:



Channel Mixing

several input channels c (RGB) and several output channels f (different feature maps):

$$o_{f,x,y} = \sum_{c,s,t} w_{f,c,s,t} i_{c,x+s,y+t}$$

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



158

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-14

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



653

+

+

+ 1 = 798



Bias = 1

one feature map


Output

-25	466	466	475	...
295	787	798		...
				...
				...
...	...	...	...	...


[source](#)

Striding and Padding

need to define how to stride over image
stride > 1 corresponds to down-sampling
→ fewer nodes after convolutional layer



zero-padding of input to make it wider:
otherwise shrinking of representation with
each layer (depending on kernel size)
→ needed for large kernels and several layers



[source](#)

Another Ingredient: Pooling



replacing outputs of neighboring nodes with summary statistic (e.g., maximum or average value of nodes)

→ non-linear downsampling (regularization)

pooling is translation invariant: no interest in exact position of, e.g., maximum value



pooling over features learned by separate kernels (cross-channel pooling) can also learn other transformation invariances, like rotation or scale

(convolutions can detect same translated motif across entire image, but not rotated or scaled versions of it)



Putting It All Together



convolutional neural networks in short:
local connections, shared weights, pooling, many layers



[source](#)

Hierarchical Representation


increasing abstraction
→




Often Beneficial: Data Augmentation

adding different variations
of input data to training

idea: supporting the model
in maintaining desired
invariances




LeNet



from wikipedia

AlexNet





AlexNet finally started the deep learning hype. (winning the ImageNet challenge in 2012)

Rise of Deep Learning

a little bit oversimplified:


deep learning = lots of training data + parallel computation + smart algorithms

AlexNet: ImageNet (with data augmentation) + GPUs + ReLU, dropout, SGD



[source](#)

ImageNet Classification top-5 error (%)



The diagram illustrates a deep convolutional neural network (CNN) architecture for image classification. It starts with an input image of size $224 \times 224 \times 3$. The network consists of several layers:

- Convolutional Layers:**
 - Layer 1: Convolution with $112 \times 112 \times 64$ filters.
 - Layer 2: Convolution with $56 \times 56 \times 256$ filters.
 - Layer 3: Convolution with $28 \times 28 \times 512$ filters.
 - Layer 4: Convolution with $14 \times 14 \times 512$ filters.
 - Layer 5: Convolution with $7 \times 7 \times 512$ filters.
- Max Pooling:** Indicated by red boxes, it is applied after the first, second, third, and fourth convolutional layers.
- Fully Connected Layers:**
 - Layer 6: A fully connected layer with $1 \times 1 \times 4096$ units.
 - Layer 7: A second fully connected layer with $1 \times 1 \times 1000$ units.
- Softmax Layer:** The final layer for classification, represented by an orange box.

Legend:

- White box: convolution+ReLU
- Red box: max pooling
- Blue box: fully connected+ReLU
- Orange box: softmax



Skip Connections


issue: degradation of training and test errors when adding more and more layers → not due to overfitting (but reason controversial)

solution: learning of residuals by means of skip connections (resulting in combination of different paths through computational graph)


→ produces loss functions that train easier

together with batch normalization (avoiding exploding gradients), skip connections enable extremely deep networks (>1000 layers) without degradation

residual mapping (special kind of skip/shortcut connections):



loss surface:



[source](#)

Foundation Models

problem with ConvNets trained from scratch (in fact, with all ML methods): only suited for domain of training data

idea of transfer learning:

pretrain a big model (called foundation model) on a broad data set, and then use these learnings for subsequent trainings on specific (typically, narrow) data

two forms of transfer learning: feature extraction and finetuning

Feature Extraction

approach:

- get a pretrained ConvNet as foundation model
- remove final fully-connected layer (its outputs are the class scores from the pretraining task)
- pass training data (for the task at hand) through the network
- for each image, extract vector of last hidden layer's activations (e.g., 4096-dimensional for AlexNet)
- use the extracted values as features to train another model (to be adjusted to the task at hand)

(same idea as used before for SIFT features)


Finetuning

approach:

- get pretrained ConvNet as foundation model
- continue (starting with weights after pretraining) training with data for task at hand

typically, need to change architecture of final layer: different number of classes

options: finetune all weights or just some



Finetuning Scenarios

finetuning gives better adjustment to new task than feature extraction (important if new task differs a lot from pretraining task)

→ coming along with danger of overfitting


ConvNet features more generic in early layers (e.g., edges) and more specific to the data set of pretraining in later layers

→ for small finetuning data sets, typically keep weights of early layers fixed to avoid overfitting

for large finetuning data sets (less danger of overfitting), finetuning all layers can be beneficial



Object Detection

output: bounding boxes with category label and confidence




PASCAL VOC Data Set

- PASCAL Visual Object Class challenge
- widely used as benchmark for object detection and semantic segmentation
- 20 object categories such as person, sofa, sheep, car, ...
- 11530 annotated images
- available annotations: pixel-level segmentation, bounding boxes, object classes




MS COCO Data Set

- Microsoft Common Objects in Context
- images of complex everyday scenes containing common objects in their natural context
- 91 objects types
- 2.5 million annotated instances in 328k images





Performance Evaluation of Localization



images with multiple objects: number of detections dependent on used confidence threshold

Object Detection: Single Object

(Classification + Localization)



Fully
Connected:
4096 to 1000

Class Scores

Cat: 0.9
Dog: 0.05
Car: 0.01
...

Correct label:
Cat

Softmax
Loss

Multitask Loss

Vector:
4096

Fully
Connected:
4096 to 4

Box
Coordinates
(x, y, w, h)

L2 Loss

Correct box:
(x', y', w', h')

Treat localization as a
regression problem!

too complicated for multiple objects

Localization for Multiple-Object Images

idea: classify many different crops of the image as object or background
(crops: sliding window over image, iterated at multiple window sizes)



Dog? No
Cat? No
Background? Yes




Dog? Yes
Cat? No
Background? No



Dog? Yes
Cat? No
Background? No



Dog? No
Cat? Yes
Background? No




issue: too many
possible crops

→ need for region proposals


Region-Based Convolutional Neural Network (R-CNN)

Per-image computation

Per-region computation for each $r_i \in r(I)$




Fast R-CNN: Crop ConvNet Features




Faster R-CNN: Use Region Proposal Network

Per-image computation



Per-region computation for each $r_i \in r(I)$

two RPN losses combined with two detection head losses during end-to-end training




Learned proposals
Shares computation with whole-image network

Single-Stage Detectors: Drop Per-Region Computation



Input image
 $3 \times H \times W$



Divide image into grid
 7×7

Image a set of base boxes
centered at each grid cell
Here $B = 3$

Within each grid cell:


- Regress from each of the B base boxes to a final box with 5 numbers:
($dx, dy, dh, dw, confidence$)
- Predict scores for each of C classes (including background as a class)
- Looks a lot like RPN, but category-specific!

Output:
 $7 \times 7 \times (5 * B + C)$


faster, but usually worse performance

Redmon et al, "You Only Look Once:
Unified, Real-Time Object Detection", CVPR 2016
Liu et al, "SSD: Single-Shot MultiBox Detector", ECCV 2016
Lin et al, "Focal Loss for Dense Object Detection", ICCV 2017

YOLO: Real-Time Object Detection



You Only Look Once:
prediction of bounding boxes
and class probabilities in one go



$S \times S$ grid on input





Image Segmentation



Classification of Each Pixel

segmentation: no objects, just pixels



GRASS, CAT, TREE,
SKY, ...


Paired training data: for each training image, each pixel is labeled with a semantic category.



At test time, classify each pixel of a new image.

minimize sum over classification losses
(cross-entropy loss at every output pixel)

Idea: Fully Convolutional, No Downsampling




replace flattened, fully-connected classification layers with 1×1 convolutions
→ maintain spatial relationships and enable pixel-wise classification:
conversion of feature maps into classification heat maps (one for each class)

but no downsampling means small receptive field and no hierarchical learning

Upsampling to the Rescue

typical spatial feature extraction: convolution and pooling layers → downsampling

upsampling layers to restore original image size



different options for down- (pooling, strided convolution)
and upsampling ...

Reverse Pooling

resampling (no learned parameters)

Nearest Neighbor

1	2
3	4




1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

unpooling (recording max positions from pooling)



Max Pooling
Remember which element was max!

1	2	6	3
3	5	2	1
1	2	2	1
7	3	4	8

Input: 4 x 4

5	6
7	8

Output: 2 x 2

Rest of the network

Max Unpooling
Use positions from pooling layer

1	2
3	4

Input: 2 x 2


0	0	2	0
0	1	0	0
0	0	0	0
3	0	0	4

Output: 4 x 4

filled with learned parameters in subsequent convolution


Reverse Convolution

convolution




not inverse convolution
→ additional learned parameters

transposed convolution



combine and sum overlaps:



[source](#)

Instance Segmentation

combination of object detection and per-instance segmentation

- object detection head: bounding box and class label
- segmentation (mask) head: binary mask for each detected object (foreground = object, background = everything else)





Image Synthesis

example: DALL-E 2

idea: generate new images as variations of training data (same distribution)

usually conditioned on text (prompt)




plenty of products: [DALL-E](#), [Stable Diffusion](#), [ImageGen](#), [Midjourney](#), ...

Stable Diffusion:



A scenic view of mountains at sunset, digital art

inpainting example ([GLIDE](#)):



“zebras roaming in the field”

prompt

Generative vs Predictive/Discriminative Models

discriminative models:

predict conditional probability $P(Y|X)$

generative models:

predict joint probability $P(Y, X)$


(or just $P(X)$ → unsupervised learning)

→ allow to generate new data samples


Generative models can be used for predictive tasks (Bayes theorem).

But predictive models are usually better at it.

discriminative model



generative model



task of generative models more difficult: need to model full data distribution rather than merely find patterns in inputs to distinguish outputs

Classic ML vs Deep Learning

Generative vs Predictive Models

text generation

```
done = False
total_reward = 0
while not done:
    state = torch.tensor(state, dtype=torch.float32)
    logits = policy_network(state)
    action = torch.argmax(logits).item()
    next_state, reward, done, _ = env.step(action)
    total_reward += reward
    state = next_state
    total_rewards.append(total_reward)

# Print average reward
print("Average Reward:", np.mean(total_rewards))
```

This code uses PyTorch to implement the policy gradient method (REINFORCE algorithm) to solve the CartPole problem. The policy network is defined as a simple feedforward neural network, and the training loop updates the policy network parameters to maximize the expected reward. Finally, it evaluates the learned policy by running it for 100 episodes and prints the average reward. Adjust hyperparameters and network architecture as needed for better performance.

Can you extend this to an actor-critic method?

Certainly! Here's an extension of the previous code using the actor-critic method to solve the CartPole problem:

Message ChatGPT

ChatGPT can make mistakes. Check important info.

ChatGPT

image synthesis



Prompt: Epic anime artwork of a wizard atop a mountain at night casting a cosmic spell into the dark sky that says "Stable Diffusion 3" made out of colorful energy

Stable Diffusion 3 — Stability AI

text-to-video



Prompt: A stylish woman walks down a Tokyo street filled with warm glowing neon and animated city signage. She wears a black leather jacket, a long red dress, and black boots, and carries a black purse. She wears sunglasses and red lipstick. She... +

Sora | OpenAI

BERT family

tabular data



dmlc XGBoost

computer vision

Classify

Detect

Segment

Track



YOLO




Deep Learning for Generative AI

Depending on the application, there are currently two dominant approaches for generative AI:

- text generation: LLMs
- image synthesis: diffusion models

note the difference between image synthesis and multimodal understanding in LLMs
(images as additional input sequences to transformer, tokenized by splitting into patches)

Generative Modeling




Objectives:

1. Learn $p_{\text{model}}(x)$ that approximates $p_{\text{data}}(x)$
2. Sampling new x from $p_{\text{model}}(x)$

Explicit density estimation: explicitly define and solve for $p_{\text{model}}(x)$

Implicit density estimation: learn model that can sample from $p_{\text{model}}(x)$ without explicitly defining it.

Different Model Types for Image Synthesis



[source](#)

→ generalization: [flow matching](#)

Practice Project

[Kaggle competition: Dogs vs. Cats](#)