

# Reinforcement Learning

Fundamentals of Artificial Intelligence

# Sequential Decision Making

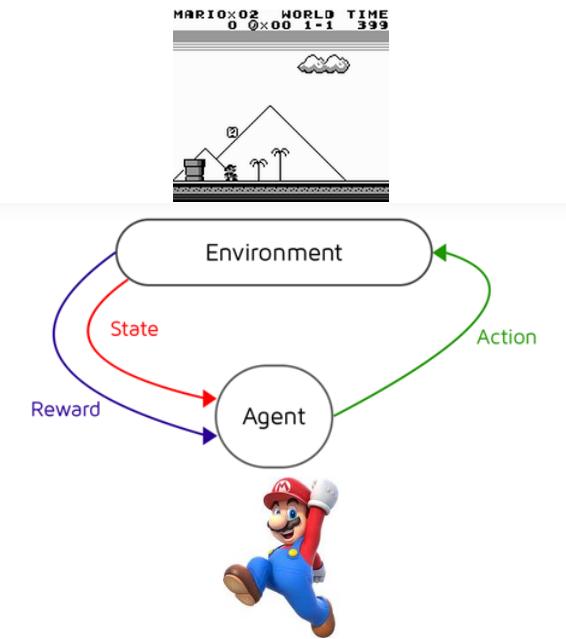
reinforcement learning (RL):

formalization of sequential decision making (action policy) of software agent interacting with environment

corresponds to search for best (or rather good) action policy to reach a given goal (e.g., win a game)

using learning from examples (data) to guide the search

RL usually more difficult (e.g., non-differentiable as a whole) than supervised learning (which can be seen as “generalized optimization”, often of proxy metric)



# Main Elements of RL

goal: find action policy maximizing reward from environment

**action policy:** exploration-exploitation trade-off

- e.g., epsilon-greedy: random exploration at small fraction of the time
- off-policy instead of on-policy learning: policy for generating observations to learn from (exploration) independent from updated policy (current best)

**feedback from environment:** goal-directed, no supervision

- scalar reward signal
- cumulative and delayed rewards (credit assignment problem)

# Optional Elements of RL


**value functions for states or actions:** improve efficiency of search in vast action policy space (alternative: direct policy search)

**model of environment:** (model-free) learning from trial-and-error or (model-based) planning

model of environment can be used in different ways:

- simulate experience from model (for learning)
- decision-time planning (e.g., heuristic search or model predictive control)

model  $\xrightarrow{\text{planning}}$  policy



from Sutton

# Markov Decision Process (MDP)


idea: current state includes all information about past

transition probabilities between states describe dynamics of given MDP

action policy: mapping from states to probabilities for selection of different actions

# States, Actions, and Rewards

transition probabilities (model of environment):  $p(s_{t+1}, r_{t+1} | s_t, a_t)$



reward hypothesis:

- reward as scalar signal
- goal: maximization of expected cumulative sum of received rewards

# State and Action Values

state/action value: total amount of expected future reward starting from given state/action (usually with discounting of later steps)

→ indicating long-term desirability of states/actions

main motivation: improve efficiency of search in policy space

(for comparison: evolutionary methods search directly by evaluating entire policies)

# State-Value Function

(needed for all states)

$$\begin{aligned} v_{\pi}(s_t) &= E_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t \right] = E_{\pi}[r_{t+1} + \gamma v_{\pi}(s_{t+1})|s_t] \\ &= \sum_{a_t} \pi(a_t | s_t) \sum_{s'_{t+1}, r_{t+1}} p(s'_{t+1}, r_{t+1} | s_t, a_t) [r_{t+1} + \gamma v_{\pi}(s'_{t+1})] \end{aligned}$$

policy: probability to take specific action being in a given state

transition probability (depending on environment) from state  $s_t$  to state  $s'_{t+1}$  for a given action

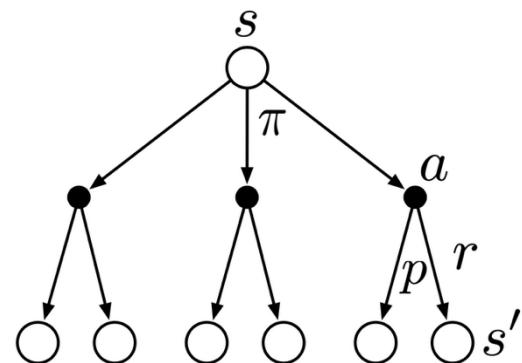
discount rate

Bellman (expectation) equation: recursion

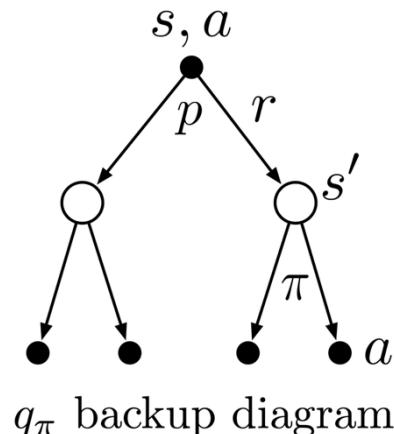
(sweep through entire state space)

# Action-Value Function

$$\begin{aligned} q_{\pi}(s_t, a_t) &= E_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t, a_t \right] = E_{\pi}[r_{t+1} + \gamma q_{\pi}(s_{t+1}, a_{t+1}) | s_t, a_t] \\ &= \sum_{s'_{t+1}, r_{t+1}} p(s'_{t+1}, r_{t+1} | s_t, a_t) \left[ r_{t+1} + \gamma \sum_{a'_{t+1}} \pi(a'_{t+1} | s'_{t+1}) q_{\pi}(s'_{t+1}, a'_{t+1}) \right] \end{aligned}$$




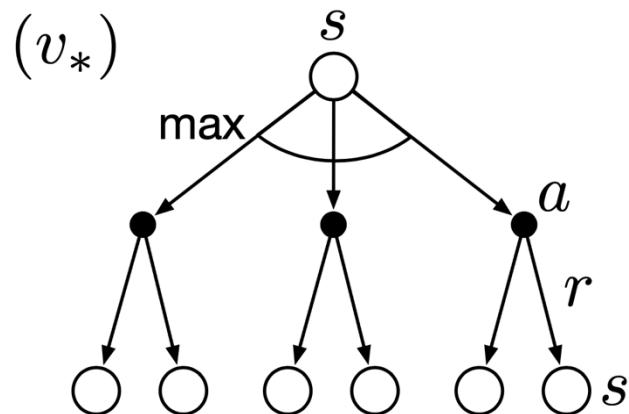
Backup diagram for  $v_{\pi}$



$q_{\pi}$  backup diagram

# Bellman Optimality Equations

optimal solutions to Bellman equations (directly defining optimal policy):



rarely possible to find in practice (due to missing model of environment, invalid Markov property, limited computational resources)  
→ approximate solutions


# Dynamic Programming (DP)

iterative approaches to find approximations for optimal value functions

1. policy evaluation: calculate value function with current policy (Bellman equation as update rule)
2. policy improvement: adjusting policy to act greedy (pick actions with maximum values) with respect to value function of current policy

putting both components together:

- policy iteration:  $\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} v_*$
- value iteration: truncated policy evaluation using Bellman optimality equation as update rule (stopped after one update of each state)



from Sutton

GPI also followed by MC and TD methods ...

# Limited Utility of DP

requires full model of environment

computationally expensive

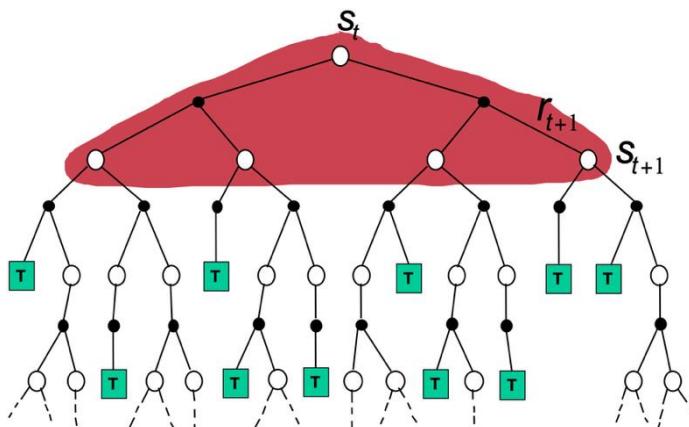
- expected update operation (based on values of all possible successor states and their probability)
- for each state (in potentially huge state space)  
(asynchronous DP at least avoids systematic sweeps over entire state space)

→ need for more efficient methods achieving the same effect as DP, without (perfect) model of environment

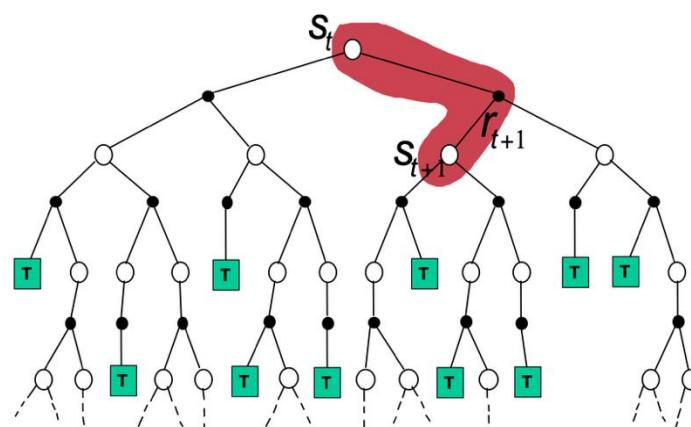
# Bootstrapping and Sampling

**bootstrapping:** update estimates of state values based on estimates of values of successor states  
**sampling:** experience of sample sequences (no need for complete knowledge of environment)


Dynamic Programming



Temporal Difference (TD) Learning



Monte Carlo (MC)



- bootstrapping
- no sampling → model-based  
(transition probabilities needed)

- bootstrapping
- sampling → model-free

- no bootstrapping
- sampling → model-free

from Sutton

# Sampling Update Rule

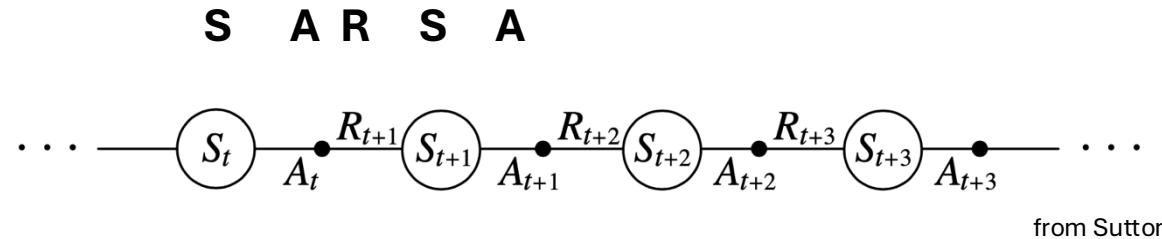
$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} [\text{Target} - \text{OldEstimate}]$$

MC:  $v(s_t) \leftarrow v(s_t) + \eta \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} - v(s_t) \right]$

TD:  $v(s_t) \leftarrow v(s_t) + \eta [r_{t+1} + \gamma v(s_{t+1}) - v(s_t)]$

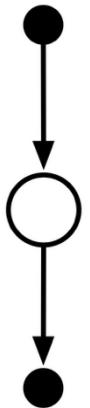
bootstrapping

# On-Policy TD Control: SARSA



following pattern of GPI:

- estimate action-value function for current behavior policy  
$$q_\pi(s_t, a_t) \leftarrow q_\pi(s_t, a_t) + \eta[r_{t+1} + \gamma q_\pi(s_{t+1}, a_{t+1}) - q_\pi(s_t, a_t)]$$
- change policy toward greediness with respect to  $q_\pi$   
(exploration for example via  $\varepsilon$ -greedy policy)

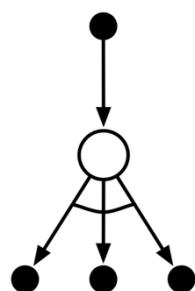


# Off-Policy TD Control: Q-Learning

estimate action-value function directly approximating optimal one  
(independent of behavior policy → potentially off-policy)


$$q(s_t, a_t) \leftarrow q(s_t, a_t) + \eta \left[ r_{t+1} + \gamma \max_a q(s_{t+1}, a_{t+1}) - q(s_t, a_t) \right]$$

policy just determines which state-action pairs are visited and updated

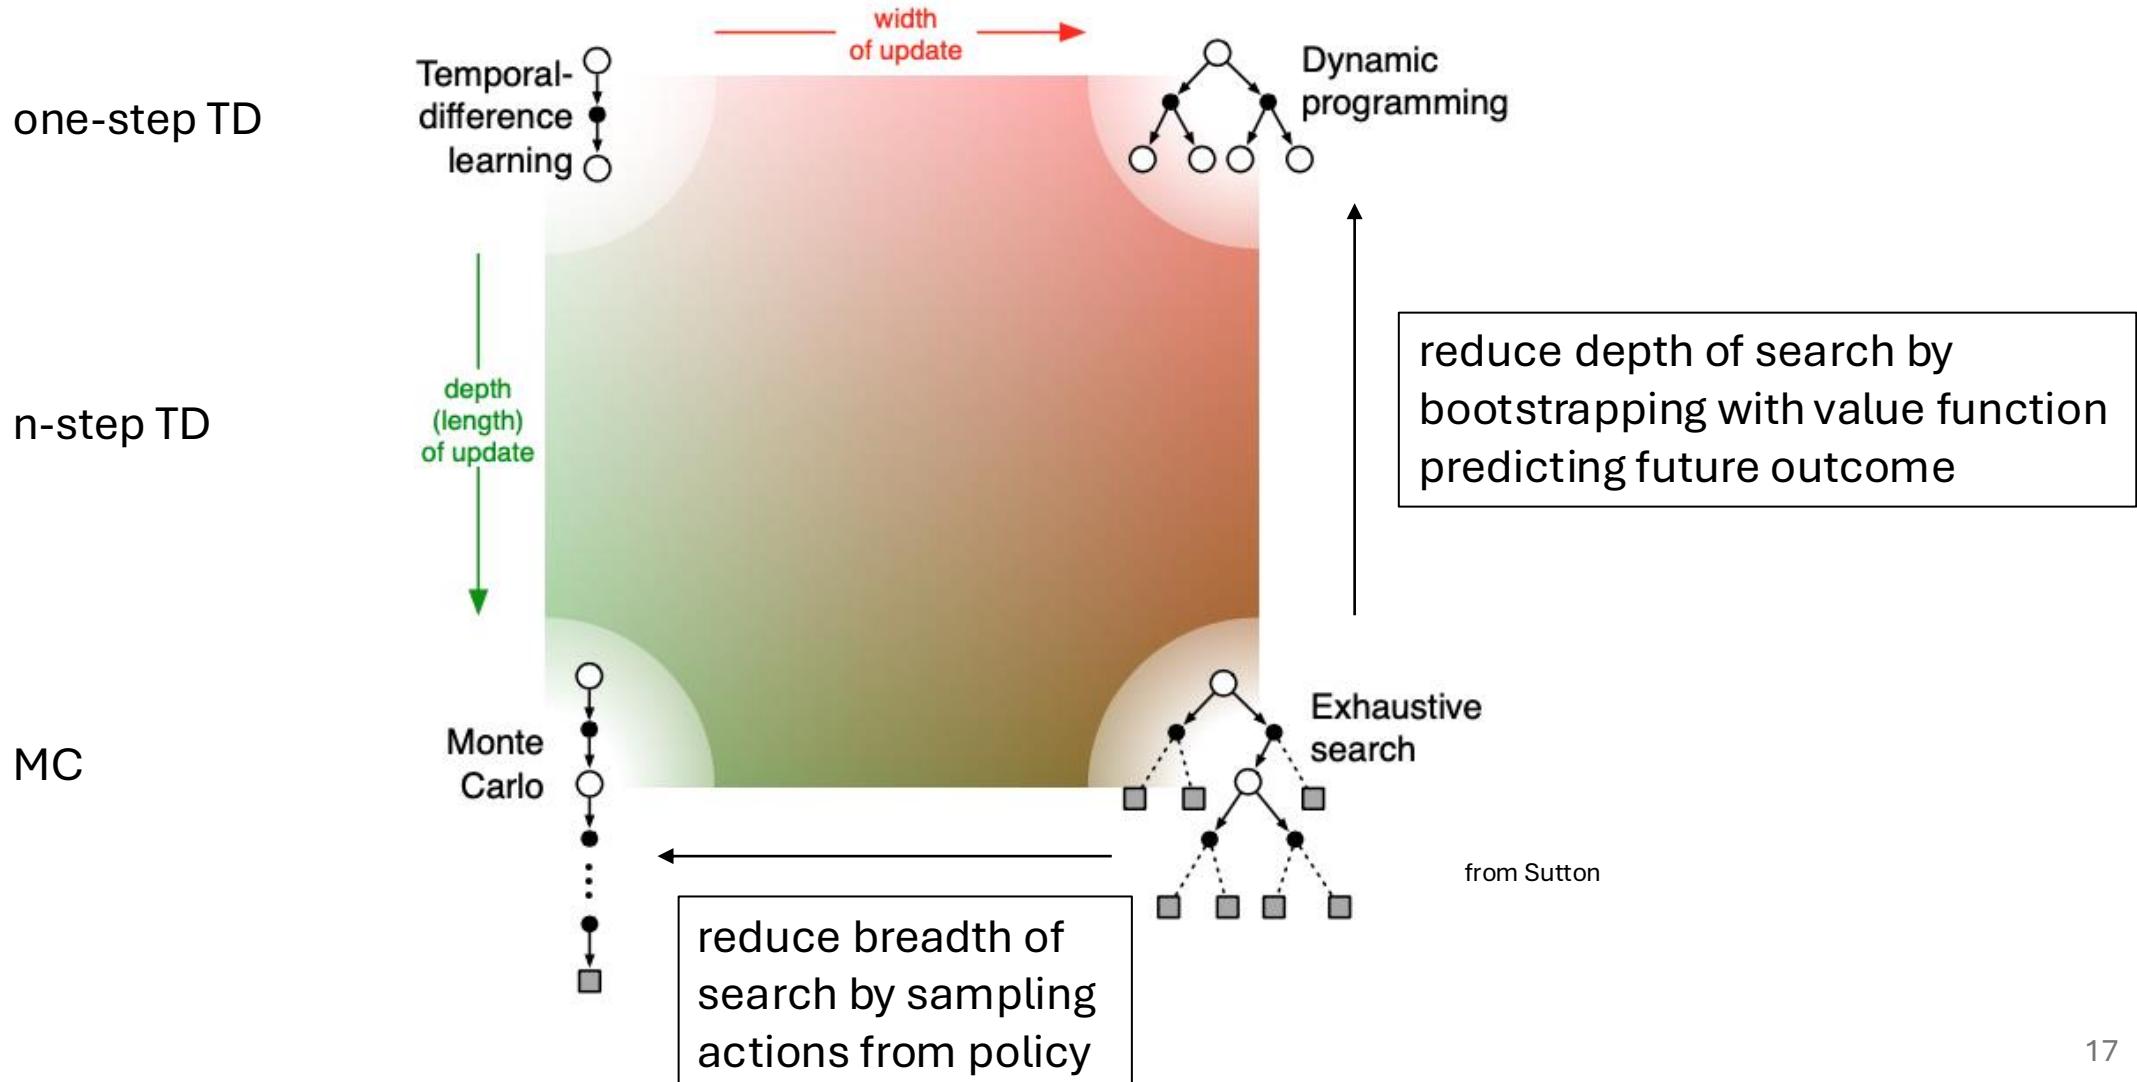


compare to expected Sarsa:

$$q(s_t, a_t) \leftarrow q(s_t, a_t) + \eta \left[ r_{t+1} + \gamma \sum_a \pi(a_{t+1}|s_{t+1}) q(s_{t+1}, a_{t+1}) - q(s_t, a_t) \right]$$



# Summary: Update Characteristics



# Limitation of Tabular Methods

tabular methods (calculating values for each state/action) simply memorize observed data

problem with tabular solution methods in practice: large state/action spaces (kind of curse of dimensionality)

- need for generalization: supervised learning to the rescue
  - non-linear function approximation over state/action space
  - nowadays often deep learning methods → deep RL

# Approximate Solution Methods

**online** setting: learning while interacting with environment  
**offline** setting: learning from limited data without further interaction with environment

state/action values as parametrized function (instead of table)

- variables/features describing different states
- parameters (e.g., connection weights in neural network) to be learned

objective function for supervised learning (e.g., squared error loss):

$$J(\hat{\mathbf{w}}) = \sum_s (\nu_{\pi}(s) - \hat{\nu}(s; \hat{\mathbf{w}}))^2$$

parameters/weights to be optimized via (stochastic) gradient descent

→ RL problem expressed in supervised learning setup (potentially offline/batch data)

but  $\nu_{\pi}(s)$  still calculated via RL methods (e.g., bootstrapping)

# Side Note: i.i.d. Assumption in ML

assumption of independent and identically distributed sets of random variables  $(Y_1, \mathbf{X}_1), (Y_2, \mathbf{X}_2), \dots, (Y_n, \mathbf{X}_n)$  fundamental to statistical (supervised) learning in terms of generalization:

consistent training and test data sets basis of empirical risk minimization  
(adversarial vulnerability/attacks: targeted violations of i.i.d. assumption)

RL: MDP outside of i.i.d. setting ( $\rightarrow$  use techniques like experience replay in training of supervised learning models for value functions with observations)

causal models: interventions outside of i.i.d. setting (need for causal model)


# The Deadly Triad

issue in deep RL: combination of off-policy bootstrapping (e.g., Q-learning) with high-dimensional function approximation leads to non-stationary targets (unstable)

most popular technique to overcome this: target networks in DQN

alternative (to conventional RL): upside-down RL

→ no bootstrapping, just supervised learning with “command” features (hindsight return in training, kind of prompt in inference)



offline RL: no interaction with environment, just fixed data set of trajectory rollouts of arbitrary policies

But policy improvement (i.e., higher return) beyond training examples (extrapolation) usually still requires policy iteration (here: iteratively updated trainings with new data with higher returns).

# Deep Q-Network (DQN)

idea: deep neural network(s) approximating tabular action-value function (according to Q-learning):  $q(s, a; \hat{w})$  as target of supervised learning model

key components to get it going:

- separate target network: weights only periodically updated with estimated Q-network weights → reducing correlations of Q-network with target (due to bootstrapping)
- experience replay: apply Q-learning updates on samples (or mini batches) of experience drawn at random from stored samples (agent's experiences) → removing correlations in observation sequence ("make it i.i.d.")

# Famous Example of Deep RL: AlphaGo

Monte Carlo tree search (heuristic, lookahead search) for move (i.e., action) selection


→ decision-time planning (advantage: focus on current state rather than full state space)

guided by deep convolutional neural networks for both value function and policy estimation

→ improving search efficiency

**reduce depth** of search tree by evaluating positions with **value function** (predicting outcome from given position → **bootstrapping**)

**reduce breath** of search tree by **sampling** actions using **policy network** (probability distribution over possible moves in given position)



# Policy Gradient Methods

learning of parametrized policy (without value functions)

$\pi(a_t|s_t; \hat{\theta})$ : probability to take different actions (target) given a state (variables/features) and parameters (e.g., neural network weights)

goal maximizing expected cumulative rewards

→ objective function corresponds to true state value:  $J(\hat{\theta}) = v_\pi(s_t)$

policy gradient theorem:

$$\nabla_{\hat{\theta}} J(\hat{\theta}) \propto \sum_{a_t} q_\pi(s_t, a_t) \nabla_{\hat{\theta}} \pi(a_t | s_t; \hat{\theta})$$

# REINFORCE


REINFORCE method (MC method following from policy gradient theorem):

$$\hat{\theta} \leftarrow \hat{\theta} + \eta \cdot \nabla_{\hat{\theta}} [\log \pi(a_t | s_t; \hat{\theta})] \cdot (r_{t+1} + \gamma r_{t+2} + \dots)$$

$$\nabla_{\hat{\theta}} J(\hat{\theta})$$

policy gradients → neural network gradients

“weighting” with observed  
(discounted) return



policy gradient methods:  
on-policy learning

# REINFORCE with Baseline

policy gradient theorem unchanged by subtracting an action-independent baseline, e.g., an estimate of the state-value function:

$$\nabla_{\hat{\theta}} J(\hat{\theta}) \propto \sum_{a_t} [q_{\pi}(s_t, a_t) - \hat{v}(s_t; \hat{w})] \nabla_{\hat{\theta}} \pi(a_t | s_t; \hat{\theta})$$

$$\hat{\theta} \leftarrow \hat{\theta} + \eta \cdot \nabla_{\hat{\theta}} [\log \pi(a_t | s_t; \hat{\theta})] \cdot [(r_{t+1} + \gamma r_{t+2} + \dots) - \hat{v}(s_t; \hat{w})]$$

e.g., separate networks

hybrid between policy-based and value-based methods  
→ reduction of variance

# Actor-Critic Methods

using state-value function for bootstrapping → critic of policy:

$$\hat{\theta} \leftarrow \hat{\theta} + \eta \cdot \nabla_{\hat{\theta}} [\log \pi(a_t | s_t; \hat{\theta})] \cdot [ (r_{t+1} + \gamma \hat{v}(s_{t+1}; \hat{w})) - \hat{v}(s_t; \hat{w}) ]$$

  
TD error

turning MC (observed return) into TD method

→ introduction of bias, but further reduction of variance

# Synonym: Advantage Actor-Critic

for the critic of the action policy (actor):

interpret TD error                     $r_{t+1} + \gamma \hat{v}(s_{t+1}; \hat{w}) - \hat{v}(s_t; \hat{w})$   
as advantage function                 $\hat{q}(s_t, a_t; \hat{w}) - \hat{v}(s_t; \hat{w})$

idea: calculates extra reward for specific action compared to average action in given state (expected state value)

Proximal Policy Optimization (PPO): prominent advantage actor-critic method with some tricks

- surrogate objective from trust region optimization → better efficiency
- clipping policy update at each training step → improved stability of actor


# Sequence Modeling for Decisions/Actions

generative: transformer decoder architecture to autoregressively model trajectories


credit assignment directly via self-attention: implicitly forming state-return associations via similarity of query and key vectors (maximizing the dot product)

desired return tokens as prompt for action generation

Decision Transformer: conditioning on desired return, past states and actions to generate future actions



Trajectory Transformer: predicting also states and returns (adding model-based components, planning with beam search)




can also condition imitation learning diffusion objective for multimodal outputs


# Robotic Control via LLMs (and Vision)

RT-2: vision-language-action model learning from web and robotics data

- representation of actions as tokens
- generalization by using pre-trained vision-language models



grounding with pre-trained skills (SayCan):



Code as Policies:

